

DS5983

PA2: The Transformer Architecture

Yunyu Guo

This report documents the implementation of two key components for a Transformer-based neural machine translation system: (1) a greedy decoding function for sentence translation (2) a comprehensive hyperparameter optimization framework. The work demonstrates practical application of the Transformer architecture for German-to-English translation with systematic performance evaluation across different model configurations.

## 1. Greedy Decoding Implementation for Sentence Translation

The `translate_sentence` function implements autoregressive text generation using greedy decoding, where the model selects the most probable next token at each step. This approach provides deterministic, fast inference suitable for real-time translation applications.

```
def translate_sentence(model, sentence, vocab_src, vocab_tgt, max_length=50):
```

### 1.2.2 Implementation Steps

Step 1: Model Preparation

```
model.eval() # Set to evaluation mode
```

Disables dropout and batch normalization training behavior

Ensures consistent inference results

Step 2: Source Preprocessing

```
src_tokens = [vocab_src['<bos>']] + [vocab_src[token] for token in tokenizer_de(sentence)] +  
[vocab_src['<eos>']]
```

```
src_tensor = torch.tensor(src_tokens, dtype=torch.long).unsqueeze(0).to(device)
```

Tokenizes German input using spaCy tokenizer

Adds beginning-of-sequence (<bos>) and end-of-sequence (<eos>) tokens

Converts to tensor format with batch dimension

Step 3: Encoder Processing

```
memory, src_mask = model.encode(src_tensor)
```

Processes source sequence through encoder stack

Returns encoded representations (memory) and attention mask

Step 4: Autoregressive Decoding

```
tgt_tokens = [vocab_tgt['<bos>']]
```

```
for _ in range(max_length):
```

```
    tgt_tensor = torch.tensor(tgt_tokens, dtype=torch.long).unsqueeze(0).to(device)
```

```
    output = model.decode(tgt_tensor, memory, src_mask)
```

```
    next_token = output[:, -1, :].argmax(-1).item()
```

```
    tgt_tokens.append(next_token)
```

```
    if next_token == vocab_tgt['<eos>']:
```

break

Greedy Selection: Chooses token with highest probability at each step

Autoregressive: Uses previously generated tokens as input

Termination: Stops when <eos> generated or max length reached

Step 5: Post-processing

```
translated_tokens = [vocab_tgt.lookup_token(token) for token in tgt_tokens[1:-1]]
```

```
translated_sentence = ''.join(translated_tokens)
```

Converts token IDs back to words

Removes special tokens (<bos>, <eos>)

Joins tokens into final sentence

### 1.3 Key Features

Deterministic Output: Greedy decoding ensures reproducible translations

Efficient Inference:  $O(n)$  time complexity for sequence length  $n$

Memory Efficient: Reuses encoder output for all decoding steps

Robust Termination: Multiple stopping criteria prevent infinite loops

### 1.4 Limitations

Limited Diversity: May miss better translations due to greedy selection

Exposure Bias: Training uses teacher forcing, but inference is autoregressive

No Backtracking: Cannot correct early poor decisions

## 2. Hyperparameter Optimization Framework

### 2.1 Overview

A systematic experimental framework was developed to evaluate the impact of key Transformer hyperparameters on translation quality. The framework enables automated training, evaluation, and comparison across multiple model configurations.

### 2.2 Experimental Design

#### 2.2.1 Hyperparameter Space

Selected Parameters:

Number of Attention Heads: [4, 8, 16]

Number of Layers: [3, 6, 9]

Learning Rate: [0.0001, 0.0005, 0.001]

Batch Size: [32, 64, 128]

Rationale:

Covers range from lightweight to computationally intensive models

Learning rates span typical optimization ranges for Transformers

Batch sizes accommodate different memory constraints

#### 2.2.2 Experimental Configuration

```
hyperparameter_configs = {  
    'num_heads': [4, 8, 16],
```

```

'num_layers': [3, 6, 9],
'learning_rate': [0.0001, 0.0005, 0.001],
'batch_size': [32, 64, 128]
}

```

# Total combinations:  $3 \times 3 \times 3 \times 3 = 81$

# Limited to 15 experiments for computational efficiency

experiments = generate\_configs(hyperparameter\_configs, max\_experiments=15)

## 2.3 Implementation Framework

### 2.3.1 Model Factory Pattern

```

def create_model_with_config(config, src_vocab_size, tgt_vocab_size, pad_idx):
    model = Transformer(
        src_vocab_size=src_vocab_size,
        tgt_vocab_size=tgt_vocab_size,
        d_model=512,                # Fixed
        N=config['num_layers'],      # Variable
        n_heads=config['num_heads'], # Variable
        d_ff=2048,                  # Fixed
        max_seq_length=5000,        # Fixed
        dropout=0.1,                # Fixed
        pad_idx=pad_idx
    )
    return model.to(device)

```

### 2.3.2 Training Protocol

Training Configuration:

Epochs per Experiment: 10 (reduced to 5 for efficiency)

Optimization: Adam optimizer with  $\beta_1=0.9$ ,  $\beta_2=0.98$

Loss Function: CrossEntropyLoss with padding token masking

Gradient Clipping: Max norm = 1.0

Evaluation: Validation loss monitoring

### 2.3.3 Metrics Collection

```

def train_with_metrics(model, train_iterator, valid_iterator, optimizer, criterion, config,
num_epochs=10):
    return {
        'config': config,
        'train_losses': train_losses,
        'val_losses': val_losses,
        'best_val_loss': best_val_loss,
        'final_train_loss': train_losses[-1],
        'final_val_loss': val_losses[-1]
    }

```

## 2.4 Experimental Results Analysis

### 2.4.1 Performance Metrics

Primary Metric: Validation loss (cross-entropy) Secondary Metrics:

Training loss convergence

Training stability

Computational efficiency

### 2.4.2 Analysis Framework

Statistical Analysis:

```
def analyze_results(results):
```

```
    # Convert to DataFrame for analysis
```

```
    # Sort by validation performance
```

```
    # Compute correlations between hyperparameters and performance
```

```
    # Identify optimal configurations
```

Visualization Components:

Bar Charts: Average performance by hyperparameter value

Training Curves: Loss progression for top configurations

Correlation Analysis: Hyperparameter impact quantification

## 2.5 Key Findings

### 2.5.1 Expected Hyperparameter Effects

Number of Attention Heads:

More heads generally improve performance up to a point

Diminishing returns beyond 8-16 heads

Computational cost increases linearly

Number of Layers:

Deeper models can capture more complex patterns

Risk of overfitting with limited training data

Training instability in very deep networks

Learning Rate:

Critical for convergence speed and final performance

Too high: unstable training, poor convergence

Too low: slow convergence, potential underfitting

Optimal range: 0.0001-0.001 for Transformers

Batch Size:

Larger batches: more stable gradients, better hardware utilization

Smaller batches: more gradient updates, potential regularization effect

Memory constraints limit upper bound

### 2.5.2 Implementation Challenges

Technical Issues Resolved:

Parameter Name Mismatch: Fixed num\_heads vs n\_heads inconsistency

Memory Management: Optimized for available GPU memory

Training Stability: Implemented gradient clipping and learning rate scheduling

### 3. System Integration and Validation

#### 3.1 End-to-End Pipeline

The complete system integrates:

Data Preprocessing: Multi30k German-English dataset

Model Training: Hyperparameter-specific configurations

Model Evaluation: Validation loss and translation quality

Translation Inference: Greedy decoding implementation

#### 3.2 Validation Approach

Translation Quality Assessment:

src\_sentence = "Ein kleiner Junge spielt draußen mit einem Ball."

translated\_sentence = translate\_sentence(model, src\_sentence, vocab\_src, vocab\_tgt)

# Expected: "A little boy playing outside with a ball."

Performance Benchmarking:

Baseline model training verification

Random data inference testing

Translation quality spot-checking

### 4. Conclusions

#### 4.1 Achievements

Successful Implementation: Both greedy decoding and hyperparameter optimization frameworks function correctly

Systematic Evaluation: Structured approach to model comparison

Practical Translation System: End-to-end German-to-English translation capability

#### 4.2 Limitations and Improvements

Current Limitations:

Greedy decoding may miss optimal translations

Limited hyperparameter exploration due to computational constraints

No advanced evaluation metrics (BLEU, METEOR)

```
.. =====
Experiment 1/15
Config: {'num_heads': 4, 'num_layers': 6, 'learning_rate': 0.0005, 'batch_size': 128}
=====
Epoch 1/10 - Train Loss: 5.7427, Val Loss: 6.8937
Epoch 2/10 - Train Loss: 5.2915, Val Loss: 11.8970
Epoch 3/10 - Train Loss: 5.2373, Val Loss: 10.9806
Epoch 4/10 - Train Loss: 5.1910, Val Loss: 11.9907
Epoch 5/10 - Train Loss: 5.1753, Val Loss: 11.1682
Epoch 6/10 - Train Loss: 5.1530, Val Loss: 11.6485
Epoch 7/10 - Train Loss: 5.1501, Val Loss: 11.4985
Epoch 8/10 - Train Loss: 5.1334, Val Loss: 11.4764
Epoch 9/10 - Train Loss: 5.1243, Val Loss: 11.5945
Epoch 10/10 - Train Loss: 5.1085, Val Loss: 12.2402

=====
Experiment 2/15
Config: {'num_heads': 4, 'num_layers': 3, 'learning_rate': 0.0005, 'batch_size': 32}
=====
Epoch 1/10 - Train Loss: 4.8733, Val Loss: 4.3700
Epoch 2/10 - Train Loss: 4.1209, Val Loss: 4.0380
Epoch 3/10 - Train Loss: 3.8970, Val Loss: 3.8139
Epoch 4/10 - Train Loss: 3.7588, Val Loss: 3.7099
Epoch 5/10 - Train Loss: 3.6813, Val Loss: 3.6829
Epoch 6/10 - Train Loss: 3.6150, Val Loss: 3.6415
Epoch 7/10 - Train Loss: 3.5662, Val Loss: 3.6067
Epoch 8/10 - Train Loss: 3.5192, Val Loss: 3.5755
Epoch 9/10 - Train Loss: 3.4726, Val Loss: 3.5741
Epoch 10/10 - Train Loss: 3.4242, Val Loss: 3.5208
```

```
=====
Experiment 3/15
Config: {'num_heads': 8, 'num_layers': 3, 'learning_rate': 0.001, 'batch_size': 128}
=====
Epoch 1/10 - Train Loss: 5.7630, Val Loss: 5.6792
Epoch 2/10 - Train Loss: 5.5046, Val Loss: 11.5638
Epoch 3/10 - Train Loss: 5.2794, Val Loss: 12.6781
Epoch 4/10 - Train Loss: 5.2329, Val Loss: 12.3284
Epoch 5/10 - Train Loss: 5.1877, Val Loss: 5.7439
Epoch 6/10 - Train Loss: 4.9255, Val Loss: 5.3508
Epoch 7/10 - Train Loss: 4.6947, Val Loss: 5.3394
Epoch 8/10 - Train Loss: 4.6441, Val Loss: 5.2400
Epoch 9/10 - Train Loss: 4.6306, Val Loss: 5.4277
Epoch 10/10 - Train Loss: 4.5485, Val Loss: 4.9637

=====
Experiment 4/15
Config: {'num_heads': 8, 'num_layers': 3, 'learning_rate': 0.0005, 'batch_size': 64}
=====
Epoch 1/10 - Train Loss: 5.0353, Val Loss: 4.6216
Epoch 2/10 - Train Loss: 4.3090, Val Loss: 4.1364
Epoch 3/10 - Train Loss: 3.9351, Val Loss: 3.8649
Epoch 4/10 - Train Loss: 3.7157, Val Loss: 3.7105
Epoch 5/10 - Train Loss: 3.5845, Val Loss: 3.6769
Epoch 6/10 - Train Loss: 3.4812, Val Loss: 3.5582
Epoch 7/10 - Train Loss: 3.4006, Val Loss: 3.4847
Epoch 8/10 - Train Loss: 3.3247, Val Loss: 3.4388
Epoch 9/10 - Train Loss: 3.2608, Val Loss: 3.4493
Epoch 10/10 - Train Loss: 3.2101, Val Loss: 3.3976
```

**Overfitting Indicators:**

Training loss decreases, validation loss increases/stagnates

Experiment 1

**Unstable Training:**

Wild validation loss fluctuations

Experiment 3 (epochs 1-4)

**Healthy Training:**

Both losses decrease together

Validation loss doesn't diverge from training

Experiments 2 and 4

Optimal Configuration (so far):

```
optimal_config = {  
    'num_heads': 4,      # Sufficient complexity  
    'num_layers': 3,     # Prevents overfitting  
    'learning_rate': 0.0005, # Stable convergence  
    'batch_size': 32     # Best generalization  
}
```